

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT on

Machine Learning (23CS6PCMAL)

Submitted by

Gopika Pushparajan (1BM23CS101)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU-560019

Jan-2026 to May-2026

B.M.S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **Gopika Pushparajan (1BM23CS101)**, who is a bonafide student of **B.M.S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Lab report has been approved as it satisfies the academic requirements in respect of an Machine Learning (23CS6PCMAL) work prescribed for the said degree.

D. Chaithra Assistant Professor Department of CSE, BMSCE	Dr. Kavitha Sooda Professor & HOD Department of CSE, BMSCE
--	--

Index

Sl. No.	Date	Experiment Title	Page No.
1	4-2-2026	Write a python program to import and export data using Pandas library functions	3
2	11-2-2026	Demonstrate various data pre-processing techniques for a given dataset	5
3	25-2-2026	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	8
4	25-3-2026	Build Logistic Regression Model for a given dataset	10
5	11-3-2026	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	13
6	25-2-2026	Build KNN Classification model for a given dataset.	16
7	1-4-2026	Build Support vector machine model for a given dataset	19
8	8-4-2026	Implement Random forest ensemble method on a given dataset.	22
9	8-4-2026	Implement Boosting ensemble method on a given dataset.	25
10	1-4-2026	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	28
11	8-4-2026	Implement Dimensionality reduction using Principal Component Analysis (PCA) method.	31


Github Link:

<https://github.com/gopikapushparajan/ML-LAB>

Program 1

Write a python program to import and export data using Pandas library functions

Screenshot



Code:

```
from sklearn.datasets import load_iris
import pandas as pd
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['species'] = iris.target
df.head()
df.columns
df.info()
df.describe()
df['species'].value_counts()

import matplotlib.pyplot as plt

plt.scatter(df['sepal length (cm)'], df['petal length (cm)'])

plt.xlabel("Sepal length")

plt.ylabel("Petal length")

plt.show()

df[df['species']==0].head()

df[df['species']==1].head()


df.corr()

df.sort_values(by='petal length (cm)').head()
```


Program 2

Demonstrate various data pre-processing techniques for a given dataset

Screenshot



Code:

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.datasets import fetch_california_housing
```

```
from sklearn.model_selection import train_test_split, GridSearchCV
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.tree import DecisionTreeRegressor

from sklearn.ensemble import RandomForestRegressor

from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

housing = fetch_california_housing(as_frame=True)

df = housing.frame

df.head()

df.shape

df.info()

plt.figure()

sns.histplot(df['MedHouseVal'], bins=30, kde=True)

plt.xlabel('Median House Value')

plt.ylabel('Frequency')

plt.title('Distribution of House Prices')

plt.show()

plt.figure()

sns.scatterplot(x=df['MedInc'], y=df['MedHouseVal'], alpha=0.5)

plt.xlabel('Median Income')

plt.ylabel('Median House Value')

plt.title('Income vs House Price')

plt.show()

plt.figure(figsize=(10,8))
```

```
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
```

```
plt.title('Feature Correlation Heatmap')
```

```
plt.show()
```

```
plt.figure()
```

```
plt.scatter(df['Longitude'], df['Latitude'],
```

```
c=df['MedHouseVal'], cmap='viridis', s=5)
```

```
plt.colorbar(label='House Value')
```

```
plt.xlabel('Longitude')
```

```
plt.ylabel('Latitude')
```




```
plt.title('House Prices by Location')
```

```
plt.show()
```


Program 3

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Screenshot:



Code:

```
import pandas as pd
```

```
from google.colab import files
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, f1_score
```

```
from sklearn.preprocessing import LabelEncoder
```

```
uploaded = files.upload()
```

```
file_path = list(uploaded.keys())[0]
```

```
df = pd.read_csv(file_path)
```

```
print("\nDataset Preview:\n", df.head())
```

```

target_column = input("\nEnter the target column name: ")

df = df.dropna()

le = LabelEncoder()

for col in df.columns:

    if df[col].dtype == 'object':

        df[col] = le.fit_transform(df[col])

X = df.drop(columns=[target_column])

y = df[target_column]

X_train, X_test, y_train, y_test = train_test_split(

    X, y, test_size=0.3, random_state=42

)

model = DecisionTreeClassifier(random_state=42)

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)

precision = precision_score(y_test, y_pred, average='weighted', zero_division=0)

recall = recall_score(y_test, y_pred, average='weighted', zero_division=0)

f1 = f1_score(y_test, y_pred, average='weighted', zero_division=0)

print("Accuracy:", accuracy)

print("\nConfusion Matrix:\n", conf_matrix)

print("\nPrecision:", precision)

print("Recall:", recall)


print("F1 Score:", f1)

```

Program 4

Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

Screenshot:



Code:

```
import pandas as pd

from google.colab import files

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

import numpy as np

import matplotlib.pyplot as plt

uploaded = files.upload()

file_path = list(uploaded.keys())[0]

df = pd.read_csv(file_path)
```

```

print("\nColumns:\n", df.columns)

print("\nPreview:\n", df.head())

feature_column = input("\nEnter ONE feature column (numeric): ")

target_column = input("Enter the target column (numeric): ")

df = df[[feature_column, target_column]].dropna()

X = df[[feature_column]]

y = df[target_column]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

model = LinearRegression()

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

mae = mean_absolute_error(y_test, y_pred)

mse = mean_squared_error(y_test, y_pred)

rmse = np.sqrt(mse)

r2 = r2_score(y_test, y_pred)

print("\n--- Simple Linear Regression Evaluation ---")

print("MAE:", mae)

print("MSE:", mse)

print("RMSE:", rmse)

print("R2 Score:", r2)


```

```
plt.figure()
plt.scatter(X_test, y_test, label="Actual Data")
plt.plot(X_test, y_pred, label="Regression Line")
plt.xlabel(feature_column)
plt.ylabel(target_column)
plt.title("Simple Linear Regression")
plt.legend()
plt.show()
```

Program 5

Build Logistic Regression Model for a given dataset

Screenshot:



Code:

```
import pandas as pd
```

```
from google.colab import files
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, f1_score
```

```
from sklearn.preprocessing import LabelEncoder, StandardScaler
```

```
uploaded = files.upload()
```

```
file_path = list(uploaded.keys())[0]
```

```
df = pd.read_csv(file_path)
```



```

print("\nDataset Preview:\n", df.head())

target_column = input("\nEnter the target column name")

df = df.dropna()

le = LabelEncoder()

for col in df.columns:

    if df[col].dtype == 'object':

        df[col] = le.fit_transform(df[col])

X = df.drop(columns=[target_column])

y = df[target_column]

X_train, X_test, y_train, y_test = train_test_split(

    X, y, test_size=0.2, random_state=42

)

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)

model = LogisticRegression(max_iter=1000)

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)

precision = precision_score(y_test, y_pred, average='weighted', zero_division=0)

recall = recall_score(y_test, y_pred, average='weighted', zero_division=0)

f1 = f1_score(y_test, y_pred, average='weighted', zero_division=0)


```

```
print("\n--- Logistic Regression Evaluation ---")  
  
print("Accuracy:", accuracy)  
  
print("\nConfusion Matrix:\n", conf_matrix)  
  
print("\nPrecision:", precision)  
  
print("Recall:", recall)  
  
print("F1 Score:", f1)
```

Program 6

Build KNN Classification model for a given dataset.

Screenshot:



Code:

```
import pandas as pd

from google.colab import files

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, f1_score

from sklearn.preprocessing import LabelEncoder, StandardScaler

uploaded = files.upload()

file_path = list(uploaded.keys())[0]

df = pd.read_csv(file_path)
```

```

print("\nDataset Preview:\n", df.head())

target_column = input("\nEnter the target column name: ")

df = df.dropna()

le = LabelEncoder()

for col in df.columns:

    if df[col].dtype == 'object':

        df[col] = le.fit_transform(df[col])

X = df.drop(columns=[target_column])

y = df[target_column]

X_train, X_test, y_train, y_test = train_test_split(

    X, y, test_size=0.2, random_state=42

)

scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)

model = KNeighborsClassifier(n_neighbors=5)

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)

precision = precision_score(y_test, y_pred, average='weighted', zero_division=0)

recall = recall_score(y_test, y_pred, average='weighted', zero_division=0)

f1 = f1_score(y_test, y_pred, average='weighted', zero_division=0)


```

```
print("\n--- KNN Evaluation ---")  
  
print("Accuracy:", accuracy)  
  
print("\nConfusion Matrix:\n", conf_matrix)  
  
print("\nPrecision:", precision)  
  
print("Recall:", recall)  
  
print("F1 Score:", f1)
```

Program 7

Build Support vector machine model for a given dataset

Screenshot:



Code:

```
import pandas as pd
```

```
from google.colab import files
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.svm import SVC
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, f1_score
```

```
from sklearn.preprocessing import LabelEncoder, StandardScaler
```

```
file_path = list(uploaded.keys())[0]
```

```
df = pd.read_csv(file_path)
```

```

print("\nDataset Preview:\n", df.head())

target_column = input("\nEnter the target column name: ")

df = df.dropna()

le = LabelEncoder()

for col in df.columns:

    if df[col].dtype == 'object':

        df[col] = le.fit_transform(df[col])

X = df.drop(columns=[target_column])

y = df[target_column]

scaler = StandardScaler()

X = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(

    X, y, test_size=0.2, random_state=42

)

model = SVC(kernel='rbf') # you can try 'linear', 'poly'

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)

precision = precision_score(y_test, y_pred, average='weighted', zero_division=0)

recall = recall_score(y_test, y_pred, average='weighted', zero_division=0)

f1 = f1_score(y_test, y_pred, average='weighted', zero_division=0)

```




```
print("\n--- SVM Evaluation ---")  
  
print("Accuracy:", accuracy)  
  
print("\nConfusion Matrix:\n", conf_matrix)  
  
print("\nPrecision:", precision)  
  
print("Recall:", recall)  
  
print("F1 Score:", f1)
```

Program 8

Implement Random forest ensemble method on a given dataset

Screenshot:



Code:

```
import pandas as pd
```

```
from google.colab import files
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, f1_score
```

```
from sklearn.preprocessing import LabelEncoder
```

```
uploaded = files.upload()
```

```
file_path = list(uploaded.keys())[0]
```

```
df = pd.read_csv(file_path)
```

```

print("\nDataset Preview:\n", df.head())

target_column = input("\nEnter the target column name: ")

df = df.dropna()

le = LabelEncoder()

for col in df.columns:

    if df[col].dtype == 'object':

        df[col] = le.fit_transform(df[col])

X = df.drop(columns=[target_column])

y = df[target_column]

X_train, X_test, y_train, y_test = train_test_split(

    X, y, test_size=0.2, random_state=42

)

model = RandomForestClassifier(

    n_estimators=100,    # number of trees

    max_depth=None,

    random_state=42

)

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)

precision = precision_score(y_test, y_pred, average='weighted', zero_division=0)

recall = recall_score(y_test, y_pred, average='weighted', zero_division=0)

```

```
f1 = f1_score(y_test, y_pred, average='weighted', zero_division=0)
print("\n--- Random Forest Evaluation ---")
print("Accuracy:", accuracy)
print("\nConfusion Matrix:\n", conf_matrix)
print("\nPrecision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

Program 9

Implement Boosting ensemble method on a given dataset.

Screenshot:

LAB - 8 : Boosting

Algorithm

1. Simple base model is trained on the entire dataset with equal weights assigned to all data points.
2. Algorithm identifies misclassified points or areas of high error.
3. A new model is added to ensemble, specifically targeting the mistakes of its predecessor.
4. All models are combined into a final predictor, often used.

Output

Accuracy = 0.93333...

Classification report:

	precision	recall	f1 score	support
0	1.00	1.00	1.00	10
1	0.89	0.89	0.89	9
2	0.91	0.91	0.91	11
accuracy			0.93	30
macro avg	0.93	0.93	0.93	30
weighted avg	0.93	0.93	0.93	30

Confusion matrix:

[10 0 0]
[0 8 1]
[0 1 10]

Feature Importance

sepal length (cm)	= 0.041886
sepal width (cm)	= 0.154977
petal length (cm)	= 0.378844
petal width (cm)	= 0.423293

Code:

```
import pandas as pd
```

```
from google.colab import files
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.ensemble import AdaBoostClassifier
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, f1_score
```

```
from sklearn.preprocessing import LabelEncoder
```

```
uploaded = files.upload()
```

```
file_path = list(uploaded.keys())[0]
```

```
df = pd.read_csv(file_path)
```

```

print("\nDataset Preview:\n", df.head())

target_column = input("\nEnter the target column name: ")

df = df.dropna()

le = LabelEncoder()

for col in df.columns:

    if df[col].dtype == 'object':

        df[col] = le.fit_transform(df[col])

X = df.drop(columns=[target_column])

y = df[target_column]

X_train, X_test, y_train, y_test = train_test_split(

    X, y, test_size=0.2, random_state=42

)

model = AdaBoostClassifier(

    n_estimators=50,

    learning_rate=1.0,

    random_state=42

)

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

conf_matrix = confusion_matrix(y_test, y_pred)

precision = precision_score(y_test, y_pred, average='weighted', zero_division=0)

```

```
recall = recall_score(y_test, y_pred, average='weighted', zero_division=0)

f1 = f1_score(y_test, y_pred, average='weighted', zero_division=0)

print("\n--- AdaBoost Evaluation ---")

print("Accuracy:", accuracy)

print("\nConfusion Matrix:\n", conf_matrix)

print("\nPrecision:", precision)



print("Recall:", recall)

print("F1 Score:", f1)
```

Program 10

Build k-Means algorithm to cluster a set of data stored in a .CSV file.

Screenshot:



Code:

```
import pandas as pd

from google.colab import files

from sklearn.cluster import KMeans

from sklearn.preprocessing import StandardScaler

import matplotlib.pyplot as plt

uploaded = files.upload()

file_path = list(uploaded.keys())[0]

df = pd.read_csv(file_path)
```



```

print("\nColumns in dataset:\n", df.columns)

print("\nDataset Preview:\n", df.head())

col1 = input("\nEnter first feature column name: ")
col2 = input("Enter second feature column name: ")

X = df[[col1, col2]].dropna()

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)

k = int(input("\nEnter number of clusters (k): "))

kmeans = KMeans(n_clusters=k, random_state=42)

y_kmeans = kmeans.fit_predict(X_scaled)

plt.figure()

for i in range(k):

    plt.scatter(

        X.iloc[y_kmeans == i, 0],

        X.iloc[y_kmeans == i, 1],

        label=f'Cluster {i}'

    )

centroids = scaler.inverse_transform(kmeans.cluster_centers_)

plt.scatter(

    centroids[:, 0],

    centroids[:, 1],

    s=200,

    c='black',

    label='Centroids'



```

```
)  
plt.xlabel(col1)  
plt.ylabel(col2)  
plt.title('K-Means Clustering')  
plt.legend()  
plt.show()
```

Program 11

Implement Dimensionality reduction using Principle Component Analysis (PCA) method

Screenshot:



Code:

```
import pandas as pd

from google.colab import files

from sklearn.preprocessing import StandardScaler

from sklearn.decomposition import PCA

import matplotlib.pyplot as plt

uploaded = files.upload()

file_path = list(uploaded.keys())[0]
```

```

df = pd.read_csv(file_path)

print("\nColumns:\n", df.columns)

print("\nPreview:\n", df.head())

use_target = input("\nDo you want to use a target column for coloring? (yes/no): ")

if use_target.lower() == 'yes':

    target_column = input("Enter target column name: ")

    y = df[target_column]

    X = df.drop(columns=[target_column])

else:

    y = None

    X = df.copy()

X = X.dropna()

X = pd.get_dummies(X, drop_first=True)

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)

pca = PCA(n_components=2)

X_pca = pca.fit_transform(X_scaled)

print("\nExplained Variance Ratio:", pca.explained_variance_ratio_)

plt.figure()

if y is not None:

    plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y)

else:

    plt.scatter(X_pca[:, 0], X_pca[:, 1])

```

```
plt.xlabel("Principal Component 1")  
plt.ylabel("Principal Component 2")  
plt.title("PCA Visualization")  
plt.show()
```